

Getting and Cleaning Data - Notes

Tanner Prestegard

Course taken from 4/6/2015 - 5/3/2015

Motivation

- Real data can be messy and incomplete, or can be in a different format than you expect.
- Data can be in an environment where you need to extract it out (ex: MySQL or MongoDB).
- The goal of this course: raw data -> processing script -> tidy data.
- After you have tidy data, you can do the data analysis (but that is not the focus of this course).

Raw and processed data

- Definition of data: data are values of qualitative or quantitative variables, belonging to a set of items.
- Raw data
 - The original source of the data.
 - Often hard to use for data analyses.
 - Data analysis includes processing.
 - Raw data may only need to be processed once.
- Processed data
 - Data that is ready for analysis
 - Processing can include merging, subsetting, transforming, etc.
 - There may be standards for processing.
 - All processing steps should be recorded!

Components of Tidy Data

- You should have four things when you finish processing data.
 - The raw data.
 - A tidy data set.
 - A code book describing each variable and its value in the tidy data set (this is also referred to as metadata).
 - An explicit and exact recipe that you used to go from raw data to the tidy data and the code book (in our case, it will be an R script).
- Raw data is in the “right” (unadulterated) format if:
 - You ran no software on the data.

- You did not manipulate or remove any of the data.
- You did not summarize the data in any way.
- Tidy data description
 - Each variable that you measured should be in one column.
 - Each different observation should be in a different row.
 - There should be one table for each “kind” of variable.
 - If you have multiple tables, they should include a column in the tables that allows them to be linked.
 - Include a row at the top of each file with variable names. These variable names should be human readable.
 - In general, data should be saved in one file per table.
- The code book
 - Information about the variables (including units) in the data set.
 - Information about the summary choices you made.
 - Information about the experimental study design you used.
 - A common format is a Word/text file.
 - There should be a section called “study design” that has a thorough description of how you collected the data.
 - There must be a section called “code book” that describes each variable and its units.
- The instruction list
 - Ideally a computer script (an R script in our case).
 - The input for the script is the raw data.
 - The output is the processed, tidy data.
 - There are no parameters in the script. (the recipe should be exact, it shouldn’t need to be tweaked or modified by future users)
 - In some cases, it will not be possible to script every step. In that case you should provide instructions. (Ex: run the script on datasets 1, 2, 3 and then run another script to combine the results)

Downloading files

- The downloading process can be included in the processing script so that everything is automated and helps with reproducibility.
- Getting/setting your working directory:
 - `getwd()`
 - `setwd()`
 - Either can use absolute or relative paths.
- Checking for and creating directories:
 - `file.exists(“directoryName”)` will check to see if the directory exists.
 - `dir.create(“directoryName”)` will create a directory if it doesn’t exist.

- Getting data from the internet - `download.file()`.
 - Downloads a file from the internet.
 - Important parameters are `url`, `destfile`, `method`.
 - * `url`: link to file.
 - * `destfile`: filename to save the data as.
 - * `method`: how to access the data. May need to use `method="curl"`.
 - Useful for downloading tab-delimited, csv, and other files.
 - Good to keep track of the date that you downloaded the data because they may change.
 - * Use the `date()` function to get this.
- Some notes:
 - If the URL starts with *http*, `download.file()` is OK.
 - If the URL starts with *https*, `download.file()` is OK in Windows.
 - If the URL starts with *https*, you may need to use `method="curl"` on a Mac.

Reading local “flat” files

- Flat files are things like text files or csv files.
- Most common way to use this is with `read.table()`.
 - Flexible and robust but requires more parameters.
 - * Important parameters: `file`, `header`, `sep`, `row.names`, `nrows`.
 - Reads the data into RAM - big data can cause problems.
 - Related: `read.csv()`, `read.csv2()`.
 - Other useful parameters:
 - * `quote`: you can tell R whether there are any quoted values. `quote=""` means no quotes.
 - * `na.strings`: sets the character that represents a missing value.
 - * `nrows`: how many rows of the file to read.
 - * `skip`: number of lines to skip before starting to read.
 - A big problem in reading flat files is that you see quotation marks like ‘ or “ in the data. Setting `quote=""` often resolves these problems.

Reading Excel files

- May be the most widely used format for sharing data, but can be a little more difficult to handle with a scripting language like R.
- Need load the `xlsx` package. `XLConnect` can be useful too.
- Functions to use:
 - `read.xlsx()`
 - `read.xlsx2()` - this is much faster but may be slightly unstable for reading subsets of rows.
- You can read specific rows and/or columns using the `colIndex` and `rowIndex` parameters.
- `write.xlsx` will write out an Excel file (and has similar arguments).
- In general it is advised to store your data as flat files (csv or tsv) as they are easier to distribute.

Reading XML

- Extensible markup language.
- Frequently used to store structured data.
- Particularly widely used in internet applications.
- Extracting XML is the basis for most web scraping.
- Components
 - Markup - labels that give the text structure.
 - Content - the actual text of the document.
- XML tags
 - Tags correspond to general labels.
 - * Start tags - ex: `<section>`
 - * End tags - ex: `</section>`
 - * Empty tags - ex: `<line-break />`
 - Elements are specific examples of tags.
 - * `<Greeting> Hello, world </Greeting>`
 - Attributes are components of the labels
 - * ``
 - * `<step number="3"> Connect A to B. </step>`
- How to read XML files into R:
 - `library(XML)`
 - `doc <- xmlTreeParse(fileURL, useInternal=TRUE)`
 - After this, it's still a structured object, so we have to use different functions to get different elements.
 - `rootNode <- xmlRoot(doc)` to get root node.
 - Can use `names(rootNode)` to get names of root node elements.
- To directly access parts of the XML document:
 - Use double brackets: `rootNode[[1]]`, or `rootNode[[1]][[1]]` if you want to go into deeper subsets (if they exist).
- To programmatically extract parts of the file:
 - Use `xmlSApply(rootNode, xmlValue)` (`xmlValue` returns the value of the element).
- A better solution may be to use XPath.
 - Problem: it's a whole new language.
 - Get information from <http://www.stat.berkeley.edu/~statcur/Workshop2/Presentations/XML.pdf>
 - `/node` is the top-level node.
 - `//node` is a node at any level.
 - We probably only need to know a few different commands in XPath.

- XPath examples:
 - `xpathSApply(rootNode,"//name",xmlValue)` will return the names of all values.
 - `xpathSApply(rootNode,"//price",xmlValue)` will return the price nodes.
 - `xpathSApply(doc,"//li[@class='score']",xmlValue)` gets all list elements with a class equal to score.
- To parse HTML files, use `htmlTreeParse()`.

Reading JSON

- JSON = Javascript Object Notation
- Lightweight data storage.
- Common format for data from application programming interfaces (APIs).
- Similar structure to XML, but a very different syntax and format.
- Data stored as:
 - Numbers (double)
 - Strings (double quoted)
 - Boolean (true or false)
 - Array (ordered, comma-separated enclosed by `[]`)
 - Object (unordered, comma-separated collection of key:value pairs enclosed by `{}`)
- Reading data from JSON
 - `library(jsonlite)`
 - `jsonData <- fromJSON(fileURL)`
 - `names(jsonData)`
- You can also take data frames in R and turn it into a JSON data set.
 - `myjson <- toJSON(iris, pretty=TRUE)` (`pretty=TRUE` gives nice indentation)
 - We can take this and send it right back to a data table using `fromJSON()`.
 - * `iris2 <- fromJSON(myjson)`

Using data.table

- Inherits from `data.frame` - all functions that accept `data.frame` work on `data.table`.
- Written in C so it is much faster in general, especially at subsetting, grouping, and updating.
- However, requires a new syntax.
- `data.table()` takes the same arguments as `data.frame()`
 - `DT <- data.frame(x=rnorm(9), y=rep(c("a","b","c"), each=3), z=rnorm(9))`
- See all the data tables in memory: `tables()`
 - Will tell you name of the data table, number of rows, memory used, column names, and if there is a "key."

- Subsetting rows is exactly the same as for a data frame.
- If you subset with only one index, it subsets by rows, which is different from a data frame.
- You can't subset columns the same way as a data frame. (ex: `DT[,c(2,3)]`)
 - You have to use expressions to summarize the data in different ways.
 - Expressions are collections of statements enclosed in curly brackets.
 - Can pass it something like `DT[,list(mean(x),sum(z))]`
- To add a new column: `DT[,w:=z^2]`, where `w` is the new column.
- If you make a copy of a data table (`DT2 <- DT`), then modified the first table, the second one will also be modified because R doesn't actually copy it to save memory when working with big datasets.
 - Need to use the copy function to actually make a copy.
- Can perform multiple operations to create a new column.
 - Ex: `DT[,m:= {tmp <- (x+z); log2(tmp+5)}]`
 - The last line is returned as the column.
- Special variables:
 - `.N`: integer of length 1, containing the number of times that a particular group appears.
 - `DT[, .N, by=x]` returns a new column showing the number of occurrences of each element.
- Keys
 - Can be useful for subsetting.


```
* DT <- data.table(x=rep(c("a","b","c",each=100), y=rnorm(300))
* setkey(DT, x)
* DT['a']
```
 - Can be used to facilitate joins for data tables.


```
* DT1 <- data.table(x=c('a','a','b','dt1'), y=1:4)
* DT2 <- data.table(x=c('a','b','dt2'), z=5:7)
* setkey(DT1, x); setkey(DT2, x)
* merge(DT1, DT2)
```
- Can be useful for fast reading of files.
 - If you save a data table in a file, you can use `fread()` to read it - this is a lot faster (about 10 times faster) than `read.table()`.

Reading from MySQL

- Free and widely used open source database software, widely used in internet-based applications.
- Data are structured in:
 - Databases
 - Tables within databases
 - Fields within tables

- Each row is called a record.
- Documentation at <http://www.mysql.com>
- Installing MySQL - go to <http://dev.mysql.com/doc/remam/5.7/en/installing.html>
- Install RMySQL on Windows:
 - Go to <http://biostat.mc.vanderbilt.edu/wiki/Main/RMySQL> or <http://www.ahschulz.de/2013/07/23/installing-rmysql-under-windows/>
- Connecting and listing databases:
 - `ucscDb <- dbConnect(MySQL(), user='genome', host = 'genome-mysql.cse.ucsc.edu')`
 - `result <- dbGetQuery(ucscDb, "show databases;"); dbDisconnect(ucscDb);`
 - result shows a list of all databases available in this MySQL server
 - To get a particular database: `hg19 <- dbConnect(MySQL(), user='genome', db='hg19', host = 'genome-mysql.cse.ucsc.edu')`
 - To get a list of all tables: `allTables <- dbListTables(hg19)`
- List field names of a table: `dbListFields(hg19, "tableName")`
- Get dimensions of a specific table: `dbGetQuery(hg19, "select count(*) from tableName")`
- To read from a table (example): `affyData <- dbReadTable(hg19, "affyU133Plus2")`
 - To select a specific subset: `query <- dbSendquery(hg19, "select * from affyU133Plus2 where misMatches between 1 and 3")`
 - `affyMis <- fetch(query); quantile(affyMis$misMatches)`
 - Can specify `n=10` as a 2nd argument in `query` to only get a few rows.
 - Note: `misMatches` is a column in the table.
 - Need to do `dbClearResult(query)` to clear the query from the remote server.
- Very important to close your connection! (do it as soon as you have the data you need)
- RMySQL vignette: <http://cran.r-project.org/web/packages/RMySQL/RMySQL.pdf>
- List of commands: <http://www.pantz.org/software/mysql/mysqlcommands.html>

Reading from HDF5

- Used for storing large data sets and structured data sets.
- Support storing a range of data types.
- HDF = Hierarchical Data Format
- Data is stored in groups which contain 0 or more groups with their metadata.
 - Have a group header with group name and list of attributes.
 - Have a group symbol table with a list of objects in the group.
- Datasets are multidimensional array of data elements along with metadata
 - Have a header with name, datatype, dataspace, and storage layout.
 - Have a data array with the data - like a data frame.

- To install and load:
 - `source("http://bioconductor.org/biocLite.R")`
 - `biocLite("rhdf5")`
 - `library(rhdf5)`
- To create a file: `created = h5createFile("example.h5")`
- To create groups within the file:
 - `created = h5createGroup("example.h5","foo")`
 - `created = h5createGroup("example.h5","foo/foobaa")`
- To list groups: `h5ls("example.h5")`
- Write to groups:
 - `A = matrix(1:10,nr=5,nc=2)`
 - `h5write(a, "example.h5", "foo/A")`
 - `B = array(seq(0.1,2.0,by=0.1),dim=c(5,2,2))`
 - `attr(B, "scale") <- "liter"`
 - `h5write(B, "example.h5", "foo/foobaa/B")`
- Write a data set directly:
 - `df = data.frame(1L:5L, seq(0,1,length.out=5), c("ab","cde","fghi","a","s"), stringAsFactors=FALSE)`
 - `h5write(df, "example.h5","df")`
- Reading data
 - `readA = h5read("example.h5", "foo/A")`
 - `readB = h5read("example.h5", "foo/foobaa/B")`
 - `readdf = h5read("example.h5", "df")`
- Writing and reading chunks
 - `h5write(c(12,13,14),"example.h5","foo/A",index=list(1:3,1))`
 - This writes to the first 3 rows in the first column of this dataset.
 - Can do the same kind of indexing with `h5read`.
- `hdf5` can be used to optimize reading/writing from disk in R.

Reading data from the web

- Web scraping is programmatically extracting data from the HTML code of websites.
 - It can be a great way to get data.
 - Many websites have information you may want to programmatically read.
 - In some cases this is against the terms of service for the website.
 - Attempting to read too many pages too quickly can get your IP address blocked.
- Getting data off webpages with `readLines()`

- `con = url("web address")`
 - `htmlCode = readLines(con)`
 - `close(con)`
- Can parse with XML instead:
 - `html = htmlTreeParse(url, useInternalNodes=TRUE)`
 - `xpathSApply(html, "//title", xmlValue)`
- Using GET from the `httr` package
 - `html2 = GET(url)`
 - `content2 = content(html2, as="test")`
 - `parsedHtml = htmlParse(content2, asText=TRUE)`
 - Then you can use `xpathSApply` as you would have otherwise.
- Accessing website with passwords
 - `pg2 = GET(url, authenticate("user", "passwd"))`
- Using handles
 - `google = handle("http://google.com")`
 - `pg1 = GET(handle=google, path="/")`
 - Only need to authenticate this handle one time.

Reading data from APIs

- API - Application Programming Interface
- Example: <https://dev.twitter.com/docs/api/1/get/blocks/blocking>
- Usually need to create an account with the development team: <https://dev.twitter.com/apps>
- Need to create a new application, fill in details.
- Accessing Twitter from R
 - `myapp = oauth_app("twitter", key="yourConsumerKey", secret="yourConsumerSecret")`
 - `sig = sign_oauth1.0(myapp, token="yourToken", token_secret="yourTokenSecret")`
 - `homeTL = GET("https://api.twitter.com/1.1/statuses/home_timeline.json", sig)`
- Converting the json object
 - `json1 = content(homeTL)`
 - `json2 = jsonlite::fromJSON(toJSON(json1))`
- How do we know what URL to use?
 - Documentation for Twitter API: <https://dev.twitter.com/docs/api/1.1/get/search/tweets>
 - Lots of other useful information there.
- `httr` allows GET, POST, PUT, DELETE requests if you are authorized.
- You can authenticate with a user name or password, but most modern APIs use something like oauth.
- `httr` works well with Facebook, Google, Twitter, Github, etc.

Reading from other sources

- There is a package in R for almost everything!
- The best way to find packages - Google “MySQL R package” for example.
- Interacting more directly with files
 - `file` - open a connection to a text file.
 - `url` - open a connection to a url.
 - `gzfile` - open a connection to a .gz file.
 - `bzfile` - open a connection to a .bz2 file.
 - `?connection` for more information.
 - Remember to close connections.
- The `foreign` packages is useful for working with data from other statistical analysis programs.
 - Loads data from Minitab, S, SAS, SPSS, Stata, Systat
 - Basic functions:
 - * `read.arff` (Weka)
 - * `read.dta` (Stata)
 - * `read.mtp` (Minitab)
 - * `read.octave` (Octave)
 - * `read.spss` (SPSS)
 - * `read.xport` (SAS)
- Reading images
 - `jpeg` - <https://cran.r-project.org/web/packages/jpeg/index.html>
 - `readbitmap` - <http://cran.r-project.org/web/packages/readbitmap/index.html>
 - `png` - <http://cran.r-project.org/web/packages/png/index.html>
- Reading GIS data (Geographic Information Systems)
 - `rdgal`, `rgeos`, `raster`
- Reading music data (directly from mp3)
 - `tuneR`, `seewave`
 - Lots of nice musical processing packages in R.

Subsetting and sorting

- Subsetting a column: `x[,1]` or `x[,"var1"]` where `var1` is the variable name.
- Subsetting rows and columns: `x[1:2,"var1"]`
- Subsetting with logical statements:
 - `x[(x$var1 <= 3 & x$var3 > 11),]`
- Dealing with missing values:
 - `x[which(x$var2 > 8),]`. The `which` command doesn't return indices of NA values.

- Sorting: `sort(x$var1, decreasing=TRUE, na.last=TRUE)`
- Ordering: sorting a data frame by a particular column's values.
 - `x[order(x$var1),]`
- Ordering with `plyr` package:
 - `library(plyr)`
 - `arrange(x, var1)` does the same thing as the `order` examples from above.
 - `arrange(x, desc(var1))` puts it in descending order.
- Adding rows and columns to data frames:
 - `x$var4 <- rnorm(5)` where `var4` was not previously a column of `x`.
 - `y <- cbind(X, rnorm(5))` does a column binding to add a column to the right side of `x`.
 - `y <- cbind(rnorm(5), x)` binds to the left side.
 - Can use `rbind` to add rows.

Summarizing data

- Look at first three rows of data frame: `head(data, n=3)`
- Look at last three rows of data frame: `tail(data, n=3)`
- To summarize: `summary(data)`.
 - Tells you minimum, median, mean, max, quartiles, etc.
- More in-depth information: `str(data)`
 - Tells you the class and size of the object (and lots of other information).
- Can look at quantiles: `quantile(data$var1, na.rm=TRUE, probs=c(0.5,0.75,0.9))`
- Make table: `table(data$var2, useNA="ifany")`
 - This tells you about the number of unique values and how often they occur.
 - If any missing values, there will be an extra table column that will tell you about the number of NAs.
 - To make a 2D table based on other variables: `table(data$var1, data$var2)`
- Check for missing values:
 - Total number of NAs: `sum(is.na(data$var1))`
 - Check if there are any NAs: `any(is.na(data$var1))`
 - Check if all elements of an array satisfies a condition: `all(data$var1 > 0)`
- Row and column sums:
 - `colSums(is.na(data))`
 - `rowSums(is.na(data))`
 - `all(colSums(is.na(data))==0)`
- Values with specific characteristics:

- `table(data$var1 %in% c("21212","21213"))`
- This returns the number of cases where this is true or false.
- Subsetting with this method: `data[data$var1 %in% c("21212","21213"),]`
- Cross tabs: identify where relationships exist in a data set.
 - `data(UCBAdmissions)`
 - `DF = as.data.frame(UCBAdmissions)`
 - `xt <- xtabs(Freq ~ Gender + Admit, data=DF)`
- Flat tables
 - `warpbreaks$replicate <- rep(1:9, len = 54)`
 - `xt = xtabs(breaks ~ ., data=warpbreaks)`
 - This creates multiple 2D tables.
 - We can summarize these tables in a larger, more compact form using `fTable(xt)`.
- Size of a data set
 - `object.size(data)` gives size in bytes.
 - `print(object.size(data), units="Mb")` gives size in MB.

Creating new variables

- Often the raw data won't have a value you are looking for, so you may need to transform the data to get what you want.
- Common variables to create:
 - “Missingness” indicators
 - “Cutting up” quantitative variables.
 - Applying transforms.
- Creating sequences (sometimes you need an index for your data set)
 - `s1 <- seq(1,10, by=2)` gives 1 3 5 7 9
 - `s2 <- seq(1,10, length=3)` gives 1.0 5.5 10.0
 - `x <- c(1,3,8,15,100); seq(along = x)` gives 1 2 3 4 5
- Subsetting variables
 - `restData$nearMe = restData$neighborhood %in% c("Roland Park", "Homeland")` adds a new subset to the data set.
- Creating binary variables
 - `restData$zipWrong = ifelse(restData$zipCode < 0, TRUE, FALSE)` gives a value of TRUE if the zip code is negative.
- Creating categorical variables out of quantitative variables
 - `restData$zipGroups = cut(restData$zipCode, breaks=quantile(restData$zipCode))`
 - Puts zip codes into four groups based on quantiles of all of the zip codes.

- Easier cutting:
 - * `library(Hmisc)`
 - * `restData$zipGroups = cut2(restData$zipCode, g=4)`
 - * Finds quantiles automatically, you don't have to set the breaks in advance.
- Cutting produces factor variables.
- Creating factor variables
 - `restData$zcf <- factor(restData$zipCode)`
- Levels of factor variables
 - `yesno <- sample(c("yes", "no"), size=10, replace=TRUE)`
 - `yesnofac = factor(yesno, levels=c("yes", "no"))`
 - `relevel(yesnofac, ref="yes")`
- Convert factor to numeric: `as.numeric(yesnofac)`
- Using the `mutate` function - creates a new version of a variable and simultaneously adds it to the data set.
 - `library(Hmisc); library(plyr)`
 - `restData2 = mutate(restData, zipGroups=cut2(zipCode, g=4))`
 - Creates a new data frame `restData2` based on `restData` and adds a new variable `zipGroups`.
- Common data transforms:
 - `abs()`, `sqrt()`, `ceiling()`, `floor()`, `round(x, digits=n)`, `signif(x, digits=n)`, `cos()`, `sin()`, `log()`, `log2()`, `log10()`, `exp()`.

Reshaping data

- Putting data into the format that you want.
- Tidy data principles (reminder)
 - Each variable forms a column.
 - Each observation forms a row
 - Each table/file stores data about one kind of observation (e.g. people/hospitals)
- Use the `reshape2` library. Examples use `mtcars` data set.
- Melting data frames
 - `mtcars$carname <- rownames(mtcars)`
 - `carMelt <- melt(mtcars, id=c("carname", "gear", "cyl"), measure.vars=c("mpg", "hp"))`
 - * We tell it which variables are measurements and which are IDs.
- Casting data frames - now we want to re-cast the data set in a particular form.
 - `cylData <- dcast(carMelt, cyl ~ variable)`
 - This gives the number of measures of each variable as a function of the number of cylinders.
 - `cylData <- dcast(carMelt, cyl ~ variable, mean)` gives the mean of each variable as a function of the number of cylinders.

- Averaging values:
 - Use `tapply(InsectSprays$count, InsectSprays$spray, mean)`
 - Another way:
 - * `spIns = split(InsectSprays$count, InsectSprays$spray)`
 - * `sprCount = lapply(spIns, mean)`
 - * Combine back into a vector: `unlist(sprCount)` or `sapply(spIns, mean)`
 - Another way: `plyr` package
 - * `ddply(Insect.Sprays,.(spray), summarize, mean=mean(count))`
- Creating a new variable:
 - `spraySums <- ddply(InsectSprays,.(spray), summarize, sum=ave(count, FUN=sum))`
- `plyr` tutorials:
 - <http://plyr.had.co.nz/0-user/>
 - <http://www.r-bloggers.com/a-quick-primer-on-split-apply-combine-problems/>
- See also the functions:
 - `acast` - for casting as multi-dimensional arrays.
 - `arrange` - for faster reordering without using `order()` commands.
 - `mutate` - for adding new variables.

Managing data frames with `dplyr`

- Data frame properties:
 - One observation per row.
 - Each column represents a variable or measure or characteristic.
- `dplyr` does not necessarily provide any new functionality, but greatly simplifies a lot of other operations and makes them much faster.
- `dplyr` function properties
 - First argument is a data frame.
 - Subsequent arguments describe what to do with it.
 - You can refer to columns in the data frame directly without using the `$` operator.
 - The result is a new data frame.
 - Data frames must be properly formatted and annotated for this to be useful.
- `select` function
 - Allows you access a set of columns within the data frame.
 - Example: `head(select(data, city:dptp))` selects all columns between those labeled `city` and `dptp`
 - `head(select(chicago,-(city:dptp)))` selects all columns except those in the specified range.
- `filter` function
 - Used to subset rows based on conditions

- `filter(chicago, pm25tmean2 > 30 & tmpd > 80)` selects all rows where these conditions are satisfied.
- **arrange** function
 - Used to re-order the rows of a data frame based on the values of a particular column.
- **rename** function
 - Used to rename a variable.
 - `chicago <- rename(chicago, pm25 = pm25tmean2, dewpoint = dptp)`
 - * New variable name is on the left!
- **mutate** function
 - Creates a new variable and adds it to the data frame.
 - `chicago <- mutate(chicago, pm25detrend = pm25-mean(pm25, na.rm = TRUE))`
- **group_by** function
 - Allows you to essentially split a data frame according to categorical variables.
 - `chicago <- mutate(chicago, tempcat = factor(1*(tmpd > 80), label=c("cold","hot"))`
 - `hotcold <- group_by(chicago, tempcat)`
 - `summarize(hotcold, pm25 = mean(pm25), o3 = max(o3tmean2), no2 = median(no2tmean2))` gives the resulting information as grouped by the “cold” and “hot” categories.
- Can use the pipeline operator `%>%` to send the output of one function directly to another function.
- `dplyr` can be useful with other data frame “backends” like `data.table` and the SQL interface for relational databases via the `DBI` package.

Merging data

- Merging more than one data set together, usually want to use some type of ID to match them up.
- Examples include two datasets: one called `reviews` and one called `solutions`.
- The `merge()` command merges data frames.
 - `x`: data frame 1
 - `y`: data frame 2
 - Use `by`, `by.x`, or `by.y` to merge.
 - Example: `merged_data = merge(reviews, solutions, by.x="solution_id", by.y="id", all=TRUE)`
 - Default - merge all common column names (can do `intersect(names(solutions), names(reviews))` to get common names).
- Can also use `join()` in the `plyr` package - faster, but fewer features.
 - Will only merge on the basis of a common ID.
 - Can use `join_all(list(df1,df2,df3))` to merge more than one data frame (must use data frame list syntax).

Editing text variables

- Fixing character vectors
 - Convert case: `tolower()`, `toupper()`
 - Splitting strings: `strsplit(str_vect,"separator")`.
 - * Example: `strsplit("Location.1","\\.")` gives "Location" "1".
 - Using `sapply()`
 - * `firstElement <- function(x){x[1]}`
 - * `sapply(splitNames,firstElement)`
 - Using `sub()`
 - * Replace underscores with nothing: `sub("_","",names(reviews))`
 - `gsub()` - replaces multiple instances. `sub()` just replaces the first occurrence.
- Finding values
 - `grep(search_str,str_vect)`, returns vector indices of matches.
 - * Can use `value=TRUE`, this returns the value rather than the index.
 - `grepl(search_str,str_vect)`, returns logical vector of TRUE or FALSE.
- Other useful functions in the `stringr` and base packages
 - `nchar(str)` tells you the number of characters in the string.
 - `substr(str,n1,n2)` gets the sub-string between the indices specified.
 - `paste(str1,str2,str3,sep=" ")` pastes the strings together using the specified separator.
 - `paste0(str1,str2,str3)` pastes together without any space.
 - `str_trim("Jeff ")` trims of all space at the end.
- Names of variables should be:
 - All lower case when possible.
 - Descriptive.
 - Not duplicated.
 - Not have underscores, dots, or white spaces.

Regular expressions

- Thought of as a combination of literals (exact matches) and metacharacters.
- Have a rich set of metacharacters.
- Used to search through strings and match patterns.
- We need a way to express whitespace word boundaries, sets of literals, the beginning/end of a line, and other things.
- Metacharacters (seem to be the same as Perl for the most part):
 - `^`: start of a line.
 - `$`: end of a line.
 - `[Bb][Uu]`: matches BU or Bu or bU or bu.

- `[0-9][a-zA-Z]`: matches any combination of number and uppercase or lowercase letters.
 - `[^?.]$:` caret here indicates “NOT”. This will match any line not ending in a question mark or period.
 - `.`: used to refer to any character. `9.10` will match `9/10`, `9110`, `911`.
 - `str1|str2|str3`: matches `str1` OR `str2` OR `str3`.
 - `?`: the indicated expression is optional.
 - * Ex: `[Gg]eorge([Ww]\.)?` `[Bb]ush` will optionally match a `W` or `w` in the middle.
 - * Note: have to escape the `.` here so it is not interpreted as a metacharacter.
 - `*`: repeats any character any number of times, including none.
 - `+`: at least one of the item.
 - `{m}`: match expression exactly `m` times.
 - `{m,}`: match expression at least `m` times.
 - `{1,5}`: match expression at least 1 time but no more than 5.
- Parentheses can also be used to return the matched expressions.
 - The matches can be accessed using `\1`, `\2`, etc.
 - Example: `+[a-zA-Z]+` `\1` + is looking for repetition of a phrase.
 - The `*` is greedy so it always matches the longest possible string that satisfies the regular expression.
 - To make it less greedy, we can use the question mark.
 - `^s(.*)s$` is less greedy than `^s(.*?)s$`.
 - Regular expressions are typically used with `grep`, `grep1`, `sub`, `gsub`.

Working with dates

- `date()` function gives a character variable containing something like “Sun Jan 12 17:48:33 2014”.
- `Sys.Date()` gives Date variable containing something like “2014-01-12”.
 - `format(d2, “%a %b %d”)` gives “Sun Jan 12”.
 - Useful characters:
 - * `%d` - day as a number.
 - * `%a` - abbreviated weekday.
 - * `%A` - unabbreviated weekday.
 - * `%m` - month (01-12).
 - * `%b` = abbreviated month.
 - * `%B` = unabbreviated month.
 - * `%y` - 2 digit year.
 - * `%Y` - 4 digit year.
- Creating dates:
 - `x <- c(“1jan1960”, “2jan1960”, “30jul1960”)`
 - `z <- as.Date(x, “%d%b%Y”)`
 - Difference: `as.numeric(z[1] - z[2])` gives -1
- `weekdays(date1)` gives day of week name as a character vector.

- `months(date1)` gives full month name as a character vector.
- `julian(date1)` gives the number of days since 1/1/1970.
 - `attr(, "origin")` gives the origin of the Julian calendar.
- lubridate packages for dates:
 - `ymd("20140108")` gives "2014-01-08 UTC"
 - `mdy("08/04/2013")` gives "2013-08-04 UTC"
 - `dmy("03-04-2013")` gives "2013-04-03 UTC"
 - Can also use things like `ymd_hms(str, tz=timezone)` on a character vector of year-month-day hour-minute-second.
 - * To learn more about how to handle timezones look at `?Sys.timezone`.
 - `wday()` gives the weekday as a number; specifying `label=TRUE` gives the day of the week abbreviation.
- `POSIXct` and `POSIXlt` are also useful date and time classes.

Data resources

- Open government sites
 - United Nations - <http://data.un.org>
 - United States - <http://www.data.gov>
 - Many more: <http://www.data.gov/opendatasites>
- Gapminder: development and human health. <http://www.gapminder.org>
- Survey data from United States: <http://www.asdfree.com>
- Infochimps Marketplace: <http://www.infochimps.com/marketplace>
- Kaggle: company that offers data science competitions. <http://www.kaggle.com>
- UCI machine learning
- Some APIs with R interfaces
 - twitter and twitterR package
 - figshare and rfigshare package
 - PLoS and rplos package
 - rOpenSci - academic focu.
 - Facebook and RFacebook
 - Google maps and RGoogleMaps